

# animal\_classification

August 15, 2024

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.models import Sequential
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.optimizers import Adam
```

```
[2]: train_dir=r'D:\ml\Animal\data\train'
test_dir=r'D:\ml\Animal\data\test'
validation_data_dir=r'D:\ml\Animal\data\validation'
```

```
[3]: train_para=ImageDataGenerator(rescale=1./255,rotation_range=40,
                                width_shift_range=0.2,
                                height_shift_range=0.2,
                                shear_range=0.2,
                                zoom_range=0.2,
                                horizontal_flip=True,
                                fill_mode='nearest')
```

```
[4]: train_generator=train_para.flow_from_directory(train_dir,
                                                    target_size=(300,300),
                                                    batch_size=50,
                                                    class_mode='categorical',
                                                    subset='training')
```

Found 13412 images belonging to 6 classes.

```
[5]: val_data = train_para.flow_from_directory(
    validation_data_dir,
    target_size=(300, 300),
    batch_size=50,
    class_mode='categorical',
    shuffle=False)
```

Found 2549 images belonging to 6 classes.

```
[6]: model=Sequential([
    Conv2D(32,(3,3), activation='relu', input_shape=(300,300,3)),
    MaxPooling2D((2,2)),

    Conv2D(64,(3,3), activation='relu', input_shape=(300,300,3)),
    MaxPooling2D((2,2)),

    Conv2D(128,(3,3), activation='relu', input_shape=(300,300,3)),
    MaxPooling2D((2,2)),

    Conv2D(128,(3,3), activation='relu', input_shape=(300,300,3)),
    MaxPooling2D((2,2)),

    Flatten(),
    Dense(512, activation='relu'),
    Dropout(0.5),
    Dense(train_generator.num_classes, activation='softmax')
])
```

```
[7]: early_stopping=EarlyStopping(monitor='val_loss',
    ↪patience=5,verbose=2,restore_best_weights=True)
```

```
[8]: model.compile(

    optimizer=Adam(learning_rate=0.001),

    loss='categorical_crossentropy',

    metrics=['Accuracy']

)
```

```
[9]: history=model.fit(train_generator,
    epochs=100,
    batch_size=64,
    validation_data=val_data,
    callbacks=[early_stopping])
```

Epoch 1/100

269/269 [=====] - 646s 2s/step - loss: 1.6042 -  
Accuracy: 0.3469 - val\_loss: 1.3375 - val\_Accuracy: 0.4845

Epoch 2/100

269/269 [=====] - 615s 2s/step - loss: 1.2729 -  
Accuracy: 0.5194 - val\_loss: 1.0661 - val\_Accuracy: 0.6034

Epoch 3/100

269/269 [=====] - 610s 2s/step - loss: 1.0687 -

Accuracy: 0.6118 - val\_loss: 0.9452 - val\_Accuracy: 0.6548  
 Epoch 4/100  
 269/269 [=====] - 606s 2s/step - loss: 0.9734 -  
 Accuracy: 0.6439 - val\_loss: 0.9467 - val\_Accuracy: 0.6536  
 Epoch 5/100  
 269/269 [=====] - 613s 2s/step - loss: 0.9045 -  
 Accuracy: 0.6770 - val\_loss: 0.8642 - val\_Accuracy: 0.6881  
 Epoch 6/100  
 269/269 [=====] - 612s 2s/step - loss: 0.8492 -  
 Accuracy: 0.6988 - val\_loss: 0.8523 - val\_Accuracy: 0.6960  
 Epoch 7/100  
 269/269 [=====] - 608s 2s/step - loss: 0.8111 -  
 Accuracy: 0.7088 - val\_loss: 0.7921 - val\_Accuracy: 0.7199  
 Epoch 8/100  
 269/269 [=====] - 615s 2s/step - loss: 0.7641 -  
 Accuracy: 0.7302 - val\_loss: 0.7309 - val\_Accuracy: 0.7481  
 Epoch 9/100  
 269/269 [=====] - 615s 2s/step - loss: 0.7406 -  
 Accuracy: 0.7369 - val\_loss: 0.7961 - val\_Accuracy: 0.7191  
 Epoch 10/100  
 269/269 [=====] - 609s 2s/step - loss: 0.7150 -  
 Accuracy: 0.7443 - val\_loss: 0.6875 - val\_Accuracy: 0.7662  
 Epoch 11/100  
 269/269 [=====] - 612s 2s/step - loss: 0.6920 -  
 Accuracy: 0.7592 - val\_loss: 0.6821 - val\_Accuracy: 0.7599  
 Epoch 12/100  
 269/269 [=====] - 611s 2s/step - loss: 0.6746 -  
 Accuracy: 0.7630 - val\_loss: 0.6757 - val\_Accuracy: 0.7595  
 Epoch 13/100  
 269/269 [=====] - 610s 2s/step - loss: 0.6712 -  
 Accuracy: 0.7657 - val\_loss: 0.6896 - val\_Accuracy: 0.7650  
 Epoch 14/100  
 269/269 [=====] - 609s 2s/step - loss: 0.6404 -  
 Accuracy: 0.7745 - val\_loss: 0.6098 - val\_Accuracy: 0.7901  
 Epoch 15/100  
 269/269 [=====] - 613s 2s/step - loss: 0.6317 -  
 Accuracy: 0.7768 - val\_loss: 0.6860 - val\_Accuracy: 0.7678  
 Epoch 16/100  
 269/269 [=====] - 611s 2s/step - loss: 0.6013 -  
 Accuracy: 0.7913 - val\_loss: 0.5862 - val\_Accuracy: 0.7944  
 Epoch 17/100  
 269/269 [=====] - 612s 2s/step - loss: 0.5882 -  
 Accuracy: 0.7955 - val\_loss: 0.5900 - val\_Accuracy: 0.7980  
 Epoch 18/100  
 269/269 [=====] - 612s 2s/step - loss: 0.5684 -  
 Accuracy: 0.8036 - val\_loss: 0.6788 - val\_Accuracy: 0.7591  
 Epoch 19/100  
 269/269 [=====] - 613s 2s/step - loss: 0.5731 -

```

Accuracy: 0.8027 - val_loss: 0.5888 - val_Accuracy: 0.7980
Epoch 20/100
269/269 [=====] - 614s 2s/step - loss: 0.5570 -
Accuracy: 0.8065 - val_loss: 0.5979 - val_Accuracy: 0.7999
Epoch 21/100
269/269 [=====] - 623s 2s/step - loss: 0.5465 -
Accuracy: 0.8110 - val_loss: 0.5859 - val_Accuracy: 0.8054
Epoch 22/100
269/269 [=====] - 613s 2s/step - loss: 0.5345 -
Accuracy: 0.8170 - val_loss: 0.5811 - val_Accuracy: 0.7991
Epoch 23/100
269/269 [=====] - 612s 2s/step - loss: 0.5260 -
Accuracy: 0.8204 - val_loss: 0.5856 - val_Accuracy: 0.8023
Epoch 24/100
269/269 [=====] - 616s 2s/step - loss: 0.5295 -
Accuracy: 0.8169 - val_loss: 0.6399 - val_Accuracy: 0.7882
Epoch 25/100
269/269 [=====] - 615s 2s/step - loss: 0.5072 -
Accuracy: 0.8244 - val_loss: 0.5207 - val_Accuracy: 0.8223
Epoch 26/100
269/269 [=====] - 614s 2s/step - loss: 0.5006 -
Accuracy: 0.8277 - val_loss: 0.5805 - val_Accuracy: 0.8023
Epoch 27/100
269/269 [=====] - 614s 2s/step - loss: 0.4982 -
Accuracy: 0.8313 - val_loss: 0.4985 - val_Accuracy: 0.8356
Epoch 28/100
269/269 [=====] - 614s 2s/step - loss: 0.4905 -
Accuracy: 0.8337 - val_loss: 0.5880 - val_Accuracy: 0.8031
Epoch 29/100
269/269 [=====] - 614s 2s/step - loss: 0.4865 -
Accuracy: 0.8318 - val_loss: 0.5054 - val_Accuracy: 0.8286
Epoch 30/100
269/269 [=====] - 615s 2s/step - loss: 0.4689 -
Accuracy: 0.8390 - val_loss: 0.5231 - val_Accuracy: 0.8250
Epoch 31/100
269/269 [=====] - 613s 2s/step - loss: 0.4615 -
Accuracy: 0.8435 - val_loss: 0.5382 - val_Accuracy: 0.8211
Epoch 32/100
269/269 [=====] - ETA: 0s - loss: 0.4773 - Accuracy:
0.8334Restoring model weights from the end of the best epoch: 27.
269/269 [=====] - 614s 2s/step - loss: 0.4773 -
Accuracy: 0.8334 - val_loss: 0.5015 - val_Accuracy: 0.8262
Epoch 32: early stopping

```

```
[14]: model.save(r'D:\ml\Animal\data\animal-classi.h5')
```

```
[15]: import numpy as np
def predict(model, img):
    img_array = tf.keras.utils.img_to_array(images[i].numpy())
    img_array = tf.expand_dims(img_array, 0)

    predictions = model.predict(img_array)

    predicted_class = class_names[np.argmax(predictions[0])]
    confidence = round(100*(np.max(predictions[0])), 0)
    return predicted_class, confidence
```

```
[27]: import tensorflow as tf
from tensorflow.keras.preprocessing.image import load_img, img_to_array
import numpy as np
import os
import matplotlib.pyplot as plt

model = tf.keras.models.load_model(r'D:\ml\Animal\data\animal-classi.h5')

class_labels = ['butterflies', 'chickens', 'elephants', 'horses', 'spiders', 'squirells']

def preprocess_image(image_path):
    img = load_img(image_path, target_size=(300, 300))
    img_array = img_to_array(img)
    img_array = np.expand_dims(img_array, axis=0)
    img_array = img_array / 255.0
    return img_array

test_images_dir = r'D:\ml\Animal\data\test'
test_images = []
for root, _, files in os.walk(test_images_dir):
    for file in files:
        if file.endswith(('.jpg', '.jpeg', '.png')):
            test_images.append(os.path.join(root, file))

plt.figure(figsize=(12, 12))
for i, image_path in enumerate(test_images[:25]):
    img_array = preprocess_image(image_path)
    prediction = model.predict(img_array)
    predicted_class = class_labels[np.argmax(prediction)]

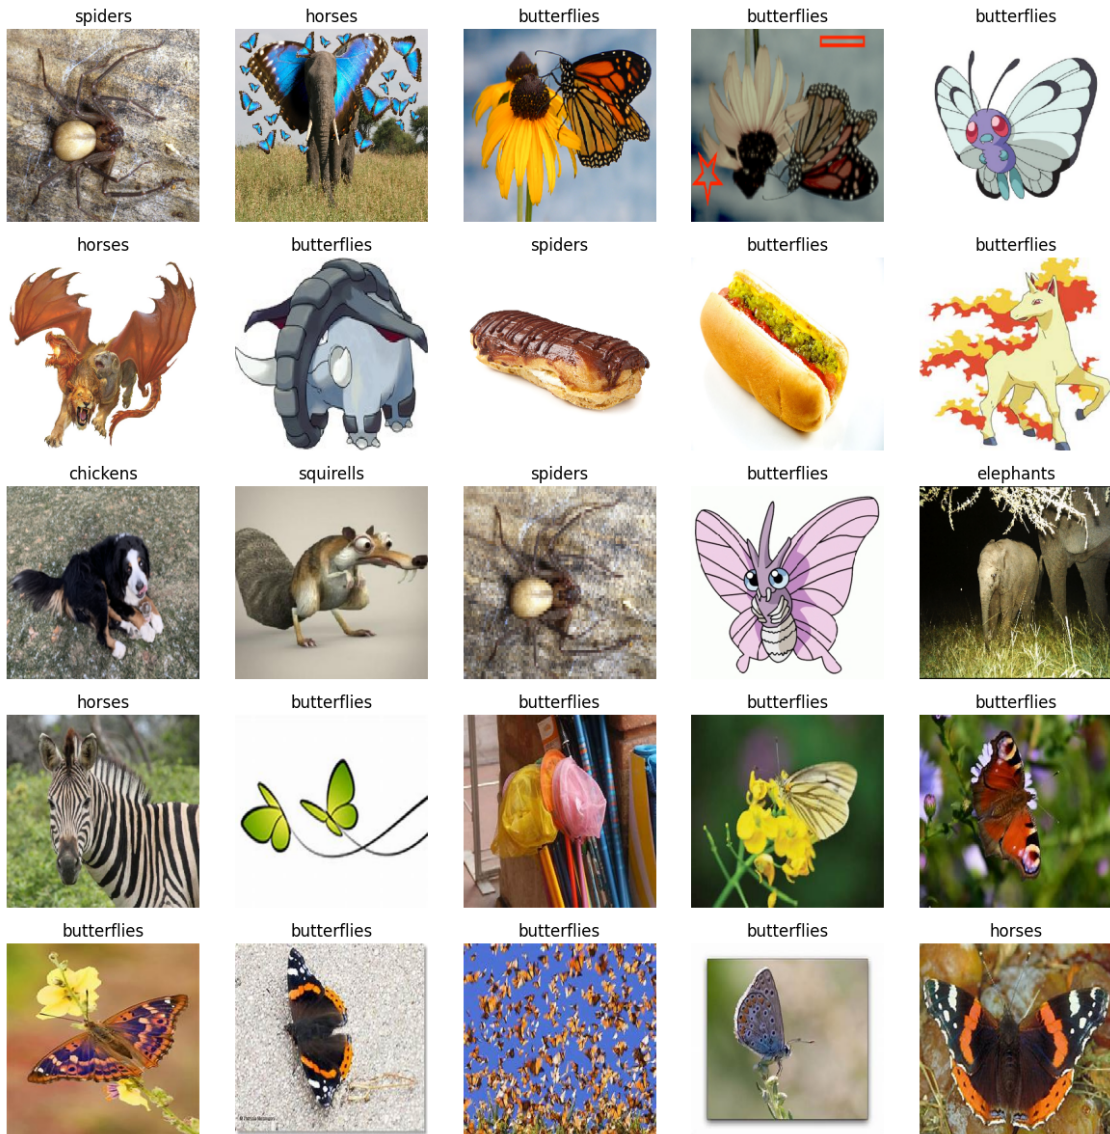
    img = load_img(image_path, target_size=(300, 300))

    plt.subplot(5, 5, i + 1)
    plt.imshow(img)
```

```
plt.title(predicted_class)
plt.axis('off')
```

```
plt.tight_layout()
plt.show()
```

```
1/1 [=====] - 0s 84ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 40ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 43ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 44ms/step
1/1 [=====] - 0s 42ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 46ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 41ms/step
1/1 [=====] - 0s 48ms/step
```



[ ]:

[ ]: